

The Lost-in-the-Middle Phenomenon in Retrieval-Augmented Generation: Mechanistic Foundations, Architectural Biases, and Mitigation Strategies

The trajectory of Large Language Model (LLM) development has been defined by an aggressive expansion of context windows, moving from the early 2,048-token limits to modern architectures capable of ingesting 128,000, one million, or even two million tokens.¹ This expansion is primarily driven by the demands of Retrieval-Augmented Generation (RAG), where external knowledge is dynamically injected into the model's prompt to ground its outputs in factual evidence.⁴ However, the capability to ingest vast volumes of data has not scaled linearly with the capability to utilize that data effectively. A critical divergence has emerged between the theoretical context window—the maximum number of tokens a model can accept—and the effective context window—the range within which a model can robustly retrieve and reason over information.⁷ This divergence is most prominently characterized by the lost-in-the-middle phenomenon, a systematic degradation in performance when the information required for a task is positioned within the central portion of a long input sequence.⁷

Theoretical Foundations and the Emergence of the U-Shaped Curve

The identification of positional bias as a fundamental constraint in modern transformer architectures was solidified through the research of Nelson Liu and colleagues in 2023, later expanded in 2024.¹ Their investigation sought to understand how LLMs actually navigate the vast contexts they are increasingly permitted to access. By designing controlled experiments that varied the position of relevant information within a long prompt, the researchers observed that model accuracy follows a distinctive U-shaped performance curve.¹¹ Accuracy is highest when the crucial evidence is placed at the beginning (primacy bias) or the end (recency bias) of the context window, but it drops significantly—often below the level of a closed-book baseline—when that information is buried in the middle.⁷

This U-shaped curve is not merely a statistical anomaly but a robust pattern observed across diverse model families, including both proprietary systems like GPT-3.5-Turbo and Claude, and open-source models like Llama.¹¹ In multi-document question-answering tasks, where exactly one document contains the answer and others serve as distractors, the performance degradation is so severe that it challenges the core premise of long-context AI.¹ For example, when GPT-3.5-Turbo was tasked with retrieving information from the middle of a long prompt, its performance fell below its ability to answer the same question without any external context,

indicating that the presence of long-form distractors actively interferes with the model's internal reasoning or parametric knowledge retrieval.¹¹

Experimental Task	Optimal Performance Positions	Observed Failure Mode	Baseline Comparison
Multi-Document QA	Start / End of Context	Substantial drop in middle accuracy	Falls below closed-book baseline ¹¹
Key-Value Retrieval	Start / End of Context	Inability to match tokens in middle	Saturation before recall limit ⁷
Needle-In-A-Haystack	Start / End of Context	"Hidden" needle in mid-sequence	Positional sensitivity at length ²
RULER Benchmark	Extreme Boundaries	Reasoning decay as context grows	Performance gap vs context size ²

The implications for RAG systems are profound. If a retriever successfully identifies a highly relevant document but the RAG pipeline appends it to the middle of a 20-document context, the generative model may effectively ignore it.¹⁵ This positional sensitivity invalidates the assumption that better retrieval (higher recall) automatically leads to better generation.⁵ Instead, the utility of retrieved information is fundamentally bounded by its placement within the final input prompt.⁵

Mechanistic Drivers of Positional Bias

Understanding why transformers exhibit such a pronounced lost-in-the-middle effect requires a deep dive into the interaction between architectural constraints, positional encodings, and training objectives.¹⁰ Recent research, particularly papers under review for ICLR 2026, proposes that this behavior is not an accidental flaw but an emergent property and a rational adaptation to the information retrieval demands encountered during the model's pre-training phase.¹⁰

Architectural Constraints and Attention Dynamics

The transformer's self-attention mechanism, while powerful, suffers from a lack of inherent spatial awareness.¹² Tokens are processed in parallel, and their relationships are determined by attention weights that can become diluted as the sequence length increases.² In autoregressive, decoder-only models (the dominant architecture for modern LLMs), the causal

mask restricts attention such that each token can only attend to preceding tokens.¹⁸ This creates a fundamental asymmetry: early tokens are attended to by every subsequent token in the sequence, while later tokens have fewer peers observing them.²⁰

A critical phenomenon in this context is the formation of attention sinks.¹⁰ Initial tokens in a sequence, such as the <s> or `` tokens, often attract disproportionately high attention weights across many layers and heads, even when they carry minimal semantic information.¹⁹ These sinks act as stable anchors for the attention mechanism, allowing the model to offload "unused" attention.¹⁹ Research indicates that disrupting these attention sinks eliminates the primacy effect but simultaneously impairs the model's ability to perform long-term memory tasks, suggesting that the primacy bias is inextricably linked to the model's structural stability.¹⁹

Positional Encoding and Long-Term Decay

The method by which models represent the position of tokens also contributes to the lost-in-the-middle effect.⁴ Many modern LLMs utilize Rotary Position Embedding (RoPE), which encodes relative positions through a series of rotations in a complex vector space.⁴ While RoPE allows for better extrapolation to longer sequences than absolute positional encodings, it often introduces a long-term decay effect.⁴ This decay causes the attention mechanism to prioritize local context (recency) and stable anchors (primacy), while the "middle" range becomes a blurred representation with lower distinctiveness.⁴ Consequently, models often fail the document metric—identifying *where* a piece of information is located—even if they can successfully perform variable extraction—identifying *what* the information is—pointing to a lack of precise spatial awareness within their own context.²¹

Adaptations to Memory Demands in Training

A compelling evolutionary perspective on the lost-in-the-middle phenomenon suggests it mirrors serial position effects in human memory.¹⁹ Human cognitive psychology has long documented the U-shaped curve in free-recall tasks, where people preferentially remember items from the beginning and end of a list.¹⁹ Researchers have demonstrated that if an LLM (such as a GPT-2 or Llama variant) is trained from scratch on two specific tasks—free recall (inducing long-term memory demand) and running span (inducing short-term memory demand)—the lost-in-the-middle behavior emerges as a natural optimization.¹⁹

Memory Demand Type	Inducing Task	Resulting Positional Effect	Architectural Requirement
Short-Term Memory	Running Span (recall last N items)	Recency Bias (end of sequence)	High local attention weights ¹⁹

Long-Term Memory	Free Recall (recall any item)	Primacy Bias (start of sequence)	Autoregressive / Attention Sinks ¹⁹
Joint Optimization	Combined Recall Tasks	U-shaped Curve (Lost-in-the-Middle)	Hybridized attention allocation ¹⁹

The recency effect directly aligns with the short-term memory demands of next-token prediction, which is the core objective of LLM pre-training.¹⁰ The primacy effect, however, is induced by uniform long-term memory demands but is distorted by the model's autoregressive architecture and the reliance on attention sinks.¹⁰ This suggests that the lost-in-the-middle effect is not a "flaw" to be "fixed" in a traditional sense, but a rational response to the way information is structured and prioritized during learning.¹⁹

Empirical Insights from the TREC 2024 RAG Track

The **TREC 2024 RAG track** provides a unique testing ground for the "lost-in-the-middle" effect because its core task involves generating **long-form, citation-backed summaries** based on realistic web search queries derived from Bing logs. Unlike synthetic "needle-in-a-haystack" tests, the track requires models to synthesize information across multiple document segments from the MS MARCO V2.1 collection.

Refuting Positional Impact in Synthesis Tasks

A significant finding emerged from the **Laboratory for Analytic Sciences (NCSU-LAS)** team. In their report, they conducted direct tests on document ordering and concluded that **"once the set of retrieved documents is fixed, reordering those documents does not affect the quality of the final summary"**. This suggests that for complex, multi-faceted synthesis tasks, the *presence* of information within the context is a more critical determinant of performance than its specific physical position, contrasting with the sharp positional sensitivity observed in simple factoid retrieval tasks.

Mitigation through Diversification and Nuggetization

Other participants explicitly designed pipelines to mitigate the phenomenon:

- **CIR Team (TH Köln):** This team hypothesized that diversified input could address the lost-in-the-middle problem. They utilized **Maximum Marginal Relevance (MMR)** to rerank segments, ensuring a variety of viewpoints were presented at the context boundaries, which prevents "information fatigue" from redundant top-ranked documents in the middle.
- **GINGER (Grounded Information Nugget-Based Generation):** Researchers found that redundant information and overly long contexts typically trigger middle loss. By operating on **"information nuggets"**—minimal, atomic units of facts—and processing them through a multi-stage pipeline (detection, clustering, ranking), they demonstrated that granular

fact-based extraction mitigates the information loss to which LLMs are prone.

- **Neural-Hungarian Optimization:** A hybrid explainability framework was proposed to **reduce positional bias** by combining semantic relevance with inter-passage attention coherence. This method uses combinatorial optimization to arrange evidence sections in a way that maximizes LLM focus and provides structured generation steps.

Evaluation and Support Assessment

The track also highlighted challenges in RAG evaluation. A large-scale comparative study by **Thakur et al. (2025)** evaluated GPT-4o as an automatic judge against human assessors for "support" assessment—verifying if cited documents actually support the generated answer. The study found that while human and LLM predictions matched in 56% to 72% of cases, LLM judges can be a reliable alternative for support assessment, though they sometimes miss evaluating information at the end of sentences or confuse similar themes.

Positional Bias in RAG Systems: The Hutter et al. Contribution

While foundational papers established the existence of the effect, subsequent research has focused on its specific manifestations within RAG pipelines.⁵ In RAG systems, the LLM does not consume information sequentially as a human reader might; rather, it processes all retrieved documents as a cohesive block within the input context.⁵ This architectural difference means that traditional Information Retrieval (IR) metrics—which assume a human-like sequential discount of information value—do not accurately predict RAG performance.⁵

Hutter and colleagues (2025) introduced the paper "Lost but Not Only in the Middle," which reveals that positional bias patterns are more complex than a simple U-shape.⁵ Different LLMs exhibit distinct "positional fingerprints," with some showing higher sensitivity to the type of distractor documents surrounding the gold document.⁵ They also found that positional bias significantly impacts the negative utility of distracting documents; a distractor placed in a "spotlight" position (the start or end) is much more likely to induce hallucinations or incorrect answers than the same distractor placed in the middle.⁵

The UDCG Metric: Rethinking RAG Evaluation

To address the failure of traditional metrics, researchers proposed the Utility and Distraction-aware Cumulative Gain (UDCG).⁵ Unlike metrics that only measure retriever success (like Recall@K), UDCG accounts for the LLM's positional bias by incorporating a "positional discount" tailored to machine processing.⁵ It quantifies both the positive contribution of relevant passages and the active degradation caused by distracting ones, providing a more accurate correlation with end-to-end answer accuracy in RAG systems.⁵ This shift is necessary because a retriever that places the most relevant document at position 8 out of 10 might as well have failed entirely, given the LLM's tendency to lose information in those

central slots.¹⁶

Impact of Model Scale and Complexity

A significant finding across several papers is that the severity of the lost-in-the-middle effect is modulated by model scale.¹⁹ Evaluations using benchmarks like LongGenBench and RULER indicate that larger models generally exhibit more resilience to positional bias.⁸

Model Category	Scaling Behavior	Performance Trend at Long Context
Open-Source (e.g., Llama-3-8B)	High sensitivity	Drastic accuracy drops (up to 47.1% on GSM8K) ²⁶
Mid-Size (e.g., Llama-3.1-70B)	Moderate robustness	Significant drops at extreme lengths (128K+) ⁸
Frontier (e.g., GPT-4o, Gemini 1.5)	High robustness	Performance saturates but rarely collapses ⁸

While increasing the model size flattens the U-shaped curve—making the recall of middle tokens more uniform—the underlying mechanism of positional bias often remains.¹⁹ Even when a model like Llama-3.2 3B or Qwen-2.5 1.5B achieves high accuracy across all positions, its internal attention weights may still show the same primacy and recency spikes, suggesting that larger models compensate for their architectural biases through sheer capacity rather than by eliminating them.¹⁹

Benchmarking Long-Context Retrieval Failures

The industry has responded to the lost-in-the-middle challenge with a suite of benchmarks designed to test information aggregation and retrieval under high-noise conditions.²

The Needle-In-A-Haystack (NIAH) Paradox

The most well-known evaluation is the Needle-In-A-Haystack test, which measures a model's ability to retrieve a specific, isolated fact hidden within a massive context of unrelated text.² While many frontier models (GPT-4o, Gemini 1.5 Pro) claim "100% NIAH accuracy," researchers have found that this metric can be misleading.¹⁴

Recent empirical results reveal that the size of the "needle" (the gold context) and the nature of the query significantly impact performance.¹⁴

- **Needle Size:** Smaller needles are significantly harder to retrieve from the middle of the

haystack than larger, more redundant ones.²⁹ For instance, Gemini-2.0-Flash saw its accuracy drop by 61 points when retrieving a small needle from the end of a context window compared to the start, whereas larger needles were more position-robust.²⁹

- **Lexical Overlap:** Benchmarks like NoLiMa (No Lexical Matching) show that models often rely on surface-level keyword matching to find the needle.⁹ When questions and needles have minimal lexical overlap, requiring latent associative reasoning, even top-tier models like GPT-4o experience a collapse in retrieval performance as context lengths scale from 1K to 32K.⁹

RULER and InfiniteBench: Measuring Aggregate Failure

Benchmarks like RULER and InfiniteBench extend the analysis from simple retrieval to complex reasoning.² RULER evaluations on GPT-4o show that even when the model *can* retrieve all relevant information perfectly, its ability to use that information to solve a problem (e.g., multi-hop QA or math) degrades substantially as input length increases.⁸ This suggests that the "lost-in-the-middle" problem is not just a retrieval failure, but a secondary "reasoning failure" that occurs because the attention mechanism cannot maintain high-fidelity representation of the relationships between scattered pieces of information.⁸

Architectural and Engineering Mitigations

Given the fundamental nature of the lost-in-the-middle problem, researchers have explored both architectural changes and system-level engineering tricks to circumvent positional bias.⁴

Strategic Context Reordering

The most immediate solution for RAG developers is the strategic ordering of retrieved documents.⁴ Instead of presenting documents in the order they were retrieved (often ranked by a single similarity score), advanced pipelines implement "long context reordering" or "attention sorting".¹⁷

Positioning Strategy	Implementation Logic	Effect on Performance
Similarity-Only	Descending similarity (1, 2, 3... N)	Most relevant document often falls in "lost" middle ¹⁵
U-Shape Alignment	Most relevant at start/end, least in middle	Leverages primacy and recency biases ⁴
Diversified Reranking	Interspersing relevant chunks with summaries	Prevents "information fatigue" in central tokens ¹⁷

Implementation of these strategies in frameworks like LlamaIndex and LangChain has become standard.³² The "LongContextReorder" module specifically identifies the retrieved nodes and repositions the most vital context to the boundaries of the prompt.³³ Studies suggest this reordering can improve RAG accuracy by as much as 15 to 30 percentage points in scenarios where a high top-K retrieval is necessary.⁴

Prompt Compression and Filtering

Another mitigation involves reducing the total volume of text the model must process.¹⁷

- **LongLLMLingua:** This technique uses a smaller language model to evaluate the information density of tokens within a long prompt and prunes those that do not contribute to the final answer.³⁵ By compressing the context, the "distance" between the start and end of the sequence is effectively shortened, moving middle-positioned information closer to the spotlight positions.³⁵
- **Contextual Summarization:** Some systems implement a "DocumentCompressorPipeline," which summarizes middle documents before they are injected into the prompt.¹⁷ This maintains the semantic signal while reducing the noise that triggers positional attention loss.¹⁷

White-Box Architectural Hacks: Ms-PoE and MoICE

Researchers have also proposed modifications to the attention mechanism itself to combat positional bias without requiring full model retraining.²²

- **Multi-Scale Positional Encoding (Ms-PoE):** This is a plug-and-play approach that enhances the model's ability to handle information in the middle by using different scales of positional information, effectively recalibrating the model's spatial resolution.²²
- **Mixture of In-Context Experts (MoICE):** For models using RoPE, MoICE introduces multiple attention heads that attend to different positions using various RoPE angles.²² This allows the model to "focus" on middle positions with the same intensity usually reserved for the beginning or end of a sequence.²²

Case Study: BriefContext in Medical RAG

High-stakes domains like healthcare provide a critical testing ground for lost-in-the-middle mitigations.²⁵ The "BriefContext" framework was specifically designed to improve the robustness of medical QA where relevant evidence might be buried deep within longitudinal patient records or lengthy clinical trials.²⁵

The Map-Reduce Strategy for Context

BriefContext addresses the lost-in-the-middle problem by decomposing a long-context task into multiple short-context sub-tasks—a "map-reduce" approach for prompts.²⁵

1. **Preflight Check:** The system uses two distinct retrieval algorithms (e.g., MedCPT for dense retrieval and BM25 for sparse retrieval) to compare rankings.²⁵ If the two rankings show high discrepancy, it flags the query as being at high risk for "lost-in-the-middle" failure.²⁵
2. **Context Map:** The retrieved documents are partitioned into smaller segments.²⁵ The LLM is then prompted to extract relevant information from each partition independently, ensuring that for each "call," the information is always in a "spotlight" position.²⁵
3. **Context Reduce:** The extracted findings are collected and summarized into a final, coherent response.²⁵

Empirical results demonstrate that BriefContext significantly outperforms standard RAG baselines when the key information is placed in the middle of a document stack.²⁵ In clinical information extraction tasks, this method reduced hallucination rates and improved clinical concept mapping accuracy by ensuring that critical but subtle symptoms buried in a medical history were not overlooked by the model's attention mechanism.¹⁷

Data Representation and Spatial Awareness

A nascent but vital area of research focuses on how the *format* of data influences its "lost-in-the-middle" profile.²¹ The GM-Extract benchmark introduced a novel distinction between "semantic understanding" (retrieving *what* a variable is) and "spatial retrieval" (identifying *where* it is).²¹

Data Format	Model Perplexity	Retrieval Accuracy Trend
Structured (Dictionary/JSON)	Lower Perplexity	Higher precision but more "brittle" positional sensitivity ²¹
Natural Language (Paragraphs)	Higher Perplexity	More robust to position but lower overall semantic precision ²¹

This research suggests that models frequently fail the "Document Metric" even when they pass the "Variable Extraction Metric," pointing to a fundamental lack of spatial awareness within their context window.²¹ Simply changing the data format can dramatically impact the perplexity and accuracy of the model, suggesting that "lost-in-the-middle" is as much a linguistic problem as it is a mathematical one.²¹

Synthesis and Future Directions

The "lost-in-the-middle" phenomenon is a testament to the fact that scaling LLMs is not a

matter of simply increasing hardware memory or attention budgets.¹ It reflects a deep-seated architectural and training bias toward boundaries—a behavior that language models share with human beings.¹⁹ As context windows continue to grow toward the multi-million token mark, the field is likely to move away from "flat" prompting toward more hierarchical and agentic retrieval strategies.¹⁵

Modern RAG pipelines are evolving into "two-stage" or "multi-stage" architectures where the first stage provides wide recall and subsequent stages perform reranking, compression, and partitioning.⁴ The emergence of benchmarks like RULER and NoLiMa has forced a more honest evaluation of what "long context" actually means, shifting the goalpost from theoretical support to effective utilization.²

The ultimate resolution of the lost-in-the-middle problem will likely require a combination of architectural breakthroughs—such as dilated attention found in LongNet or state-space models like Mamba—and sophisticated engineering workflows that treat the LLM context not as a static bucket of tokens, but as a dynamic and biased attention space.¹² For current researchers and engineers, the "lost-in-the-middle" effect serves as a critical design constraint: in the world of long-context AI, position is as important as content, and the middle is where information goes to die.¹⁶

Works cited

1. From 2K to 2M+ Tokens: The Long-Context Frontiers of GenAI - DataDrivenInvestor, accessed April 16, 2026, <https://medium.datadriveninvestor.com/from-2k-to-2m-tokens-the-long-context-frontiers-of-genai-7b7304dcdabc>
2. LLM Interview Series: Context Windows, Memory, and Long-context Reasoning - Medium, accessed April 16, 2026, <https://medium.com/@huanzidage/llm-interview-series-context-windows-memory-and-long-context-reasoning-84bdb3ca0e0b>
3. Long Context Models: Working with 1M+ Token Windows - Let's Data Science, accessed April 16, 2026, <https://letsdatascience.com/blog/long-context-models-working-with-1m-token-windows>
4. Solving the 'Lost in the Middle' Problem: Advanced RAG Techniques for Long-Context LLMs, accessed April 16, 2026, <https://www.getmaxim.ai/articles/solving-the-lost-in-the-middle-problem-advanced-rag-techniques-for-long-context-llms/>
5. Lost but Not Only in the Middle: Positional Bias in Retrieval Augmented Generation | Request PDF - ResearchGate, accessed April 16, 2026, https://www.researchgate.net/publication/390507717_Lost_but_Not_Only_in_the_Middle_Positional_Bias_in_Retrieval_Augmented_Generation
6. Improving RAG accuracy: 10 techniques that actually work - Redis, accessed April 16, 2026, <https://redis.io/blog/10-techniques-to-improve-rag-accuracy/>
7. Lost in the Middle: How Language Models Use Long Contexts - MIT Press Direct,

accessed April 16, 2026,

https://direct.mit.edu/tac/article/doi/10.1162/tac/a_00638/119630/Lost-in-the-Middle-How-Language-Models-Use-Long

8. Context Length Alone Hurts LLM Performance ... - ACL Anthology, accessed April 16, 2026, <https://aclanthology.org/2025.findings-emnlp.1264.pdf>
9. NoLiMa: Long-Context Evaluation Beyond Literal Matching - arXiv, accessed April 16, 2026, <https://arxiv.org/html/2502.05167v3>
10. Lost in the Middle: An Emergent Property from Information Retrieval Demands in LLMs, accessed April 16, 2026, https://www.researchgate.net/publication/396462441_Lost_in_the_Middle_An_Emergent_Property_from_Information_Retrieval_Demands_in_LLMs
11. Lost in the Middle: How Language Models Use Long Contexts, accessed April 16, 2026, <https://cs.stanford.edu/~nfliu/papers/lost-in-the-middle.arxiv2023.pdf>
12. Lost in the Middle: How Language Models Use Long Contexts, accessed April 16, 2026, https://teapot123.github.io/files/CSE_5610_Fall25/Lecture_12_Long_Context.pdf
13. Daily Papers - Hugging Face, accessed April 16, 2026, <https://huggingface.co/papers?q=lost-in-the-middle%20scenarios>
14. NoLiMa: Long-Context Evaluation Beyond Literal Matching - arXiv, accessed April 16, 2026, <https://arxiv.org/html/2502.05167v2>
15. Lost in the Middle in LLMs. Why large language models ignore the... | by Cengizhan Bayram | Medium, accessed April 16, 2026, <https://medium.com/@cenghanbayram35/lost-in-the-middle-in-llms-86e461dc7212>
16. How to Evaluate RAG Systems: Metrics, Methods, and What to Measure First - Comet, accessed April 16, 2026, <https://www.comet.com/site/blog/rag-evaluation/>
17. Mastering the Lost in the Middle Problem in RAG | by Abheshith - Medium, accessed April 16, 2026, <https://medium.com/@abheshith7/mastering-the-lost-in-the-middle-problem-in-rag-e08482780b0f>
18. Lost but Not Only in the Middle - UvA-DARE (Digital Academic Repository), accessed April 16, 2026, https://pure.uva.nl/ws/files/246521941/978-3-031-88708-6_16.pdf
19. LOST IN THE MIDDLE: AN EMERGENT PROPERTY ... - OpenReview, accessed April 16, 2026, <https://openreview.net/pdf?id=XSHp62BCXN>
20. Lost in the Middle: An Emergent Property from Information Retrieval Demands in LLMs, accessed April 16, 2026, <https://arxiv.org/html/2510.10276v1>
21. What Works for 'Lost-in-the-Middle' in LLMs? A Study on GM-Extract and Mitigations - arXiv, accessed April 16, 2026, <https://arxiv.org/html/2511.13900v1>
22. Found in the Middle: How Language Models Use Long Contexts Better via Plug-and-Play Positional Encoding - ResearchGate, accessed April 16, 2026, https://www.researchgate.net/publication/397200235_Found_in_the_Middle_How_Language_Models_Use_Long_Contexts_Better_via_Plug-and-Play_Positional_Encoding
23. Redefining Retrieval Evaluation in the Era of LLMs - ACL Anthology, accessed April

- 16, 2026, <https://aclanthology.org/2026.eacl-long.391.pdf>
24. Redefining Retrieval Evaluation in the Era of LLMs - arXiv, accessed April 16, 2026, <https://arxiv.org/html/2510.21440v1>
25. Leveraging long context in retrieval augmented language models for ..., accessed April 16, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12048518/>
26. LongGenBench: Long-context Generation Benchmark - arXiv, accessed April 16, 2026, <https://arxiv.org/html/2410.04199v1>
27. LONGGENBENCH: Long-context Generation Benchmark - ACL Anthology, accessed April 16, 2026, <https://aclanthology.org/2024.findings-emnlp.48.pdf>
28. ICML Poster NoLiMa: Long-Context Evaluation Beyond Literal Matching, accessed April 16, 2026, <https://icml.cc/virtual/2025/poster/46685>
29. Lost in the Haystack: Smaller Needles are More Difficult for LLMs to ..., accessed April 16, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12478432/>
30. HELMET: HOW TO EVALUATE LONG-CONTEXT LLMs, accessed April 16, 2026, <https://par.nsf.gov/servlets/purl/10583917>
31. 5 Techniques for Efficient Long-Context RAG - MachineLearningMastery.com, accessed April 16, 2026, <https://machinelearningmastery.com/5-techniques-for-efficient-long-context-rag/>
32. A Practical Approach to Retrieval Augmented Generation Systems - 4 From Simple to Advanced RAG - GitHub Pages, accessed April 16, 2026, https://mallahyari.github.io/rag-ebook/04_advanced_rag.html
33. Setting up Query Pipeline For Advanced RAG Workflow using LlamaIndex - AI Planet, accessed April 16, 2026, <https://medium.aiplanet.com/setting-up-query-pipeline-for-advanced-rag-workflow-using-llamaindex-666ddd7d0d41>
34. LlamaIndex update 2023-10-10, accessed April 16, 2026, <https://www.llamaindex.ai/blog/llamaindex-update-2023-10-10-3718a3d19fb9>
35. 12 RAG Pain Points and Solutions - Origins AI, accessed April 16, 2026, <https://originshq.com/blog/12-rag-pain-points-and-solutions/>
36. Found in the Middle: How Language Models Use Long Contexts Better via Plug-and-Play Positional Encoding - arXiv, accessed April 16, 2026, <https://arxiv.org/html/2403.04797v1>
37. Contradictions in Context: Challenges for Retrieval-Augmented Generation in Healthcare, accessed April 16, 2026, <https://arxiv.org/html/2511.06668v2>
38. (PDF) Leveraging long context in retrieval augmented language models for medical question answering - ResearchGate, accessed April 16, 2026, https://www.researchgate.net/publication/391419293_Leveraging_long_context_in_retrieval_augmented_language_models_for_medical_question_answering
39. Beyond Fine-Tuning: Robust Food Entity Linking under Ontology Drift with FoodOntoRAG, accessed April 16, 2026, <https://arxiv.org/html/2603.09758v1>
40. A neural-symbolic AI agent system for biomedical concept mapping - ResearchGate, accessed April 16, 2026, https://www.researchgate.net/publication/403506725_A_neural-symbolic_AI_agent_system_for_biomedical_concept_mapping

41. Chunking, Retrieval, and Re-ranking: An Empirical Evaluation of RAG Architectures for Policy Document Question Answering - arXiv, accessed April 16, 2026, <https://arxiv.org/html/2601.15457v1>